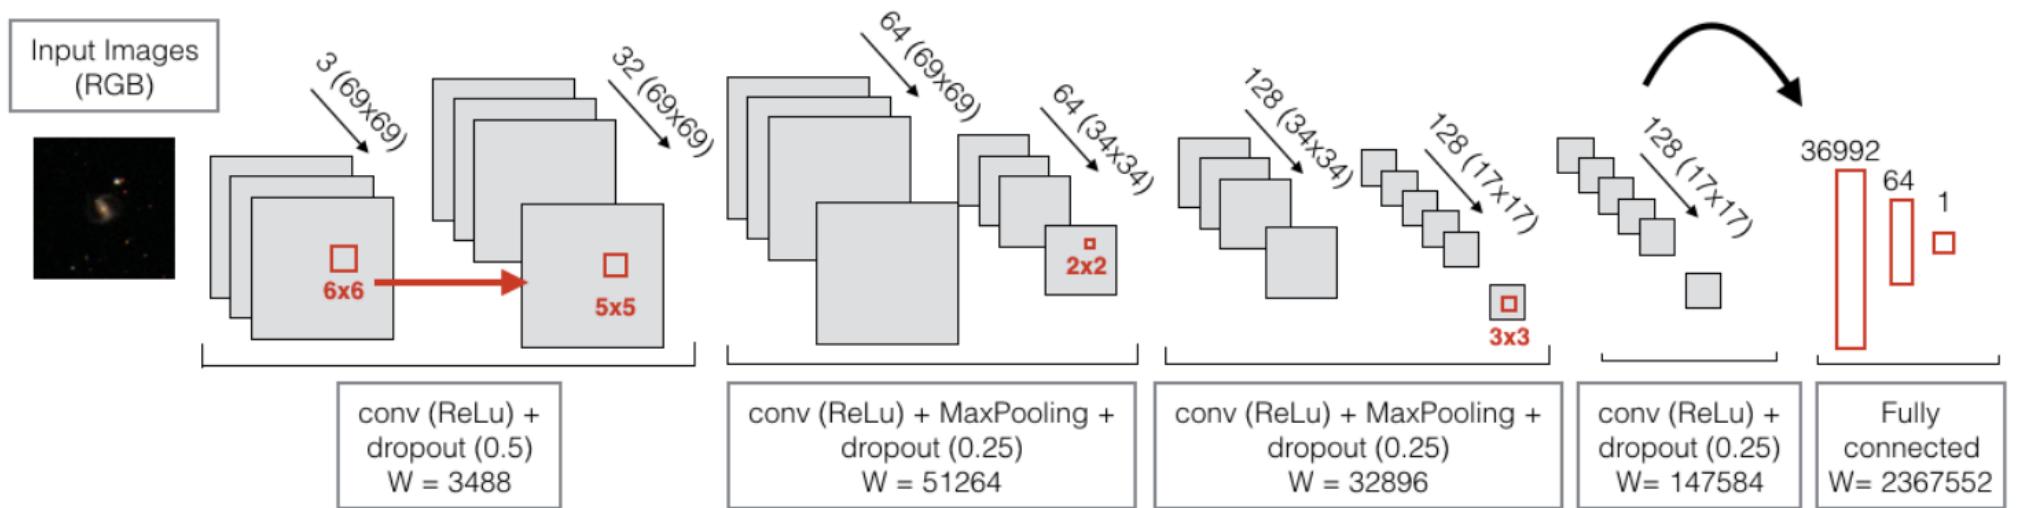


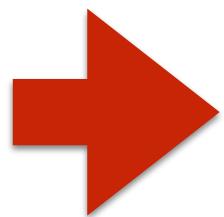
BEYOND CLASSIFICATION: IMAGE2IMAGE NETWORKS

UP TO NOW CNNs MAP IMAGES (SIGNALS) INTO FLOATS



Dominguez-Sanchez+18

Classification has its limits




HOW DO I CLASSIFY THIS IMAGE?

Classification has its limits



classification

person, sheep, dog





HOW DO I CLASSIFY THIS IMAGE?

Classification has its limits



classifi

per



HOW DO I CLASSIFY THIS IMAGE?

Going beyond classification: increasing complexity


classification



semantic segmentation



object detection



instance segmentation



Going beyond classification: increasing complexity


classification



semantic segmentation




object detection





instance segmentation



UNPOOLING OPERATION (INVERSE OF POOLING)







COPY PIXELS IN A
GIVEN WINDOW



GENERATES
LARGER IMAGES
FROM SMALLER
ONES

TRANSPOSED CONVOLUTION

ALLOWS TO INCREASE THE SIZE



Going Backward of Convolution

EXAMPLE TAKEN FROM HERE

CONVOLUTION MATRIX

	0	1	2
0	1	4	1
1	1	4	3
2	3	3	1

Kernel (3, 3)


THE KERNEL CAN BE ARRANGED IN FORM OF A MATRIX:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	1	4	1	0	1	4	3	0	3	3	1	0	0	0	0	0
1	0	1	4	1	0	1	4	3	0	3	3	1	0	0	0	0
2	0	0	0	0	1	4	1	0	1	4	3	0	3	3	1	0
3	0	0	0	0	0	1	4	1	0	1	4	3	0	3	3	1

Convolution Matrix (4, 16)


EXAMPLE TAKEN FROM HERE

THE INPUT IS FLATTENED INTO A COLUMN VECTOR




EXAMPLE TAKEN FROM HERE

THE CONVOLUTION IS TRANSFORMED INTO A PRODUCT OF MATRICES



EXAMPLE TAKEN FROM HERE


THE TRANSPOSED CONVOLUTION IS THE INVERSE OPERATION



**RESHAPED
OUTPUT**


EXAMPLE TAKEN FROM HERE

ENCODER-DECODERS GO FROM IMAGE 2 IMAGE




WE CALL THIS FULLY CONVOLUTIONAL
NEURAL NETWORKS



ENCODING-DECODING TO EXTRACT IMAGE FEATURES: U-NET



ENCODING-DECODING TO EXTRACT IMAGE FEATURES: THE U-NET



Painting Baryons



Harrington+21


Neural Networks are used to learn the non-linear mapping between cheap dark matter only simulations to expensive baryonic physics

Rodriguez+19, Modi+18, Berger+18, He+18, Zhang+19, Troster+19, Zamudio-Fernandez+19, Perraudin+19, Charnock+19, List+19, Giusarma+19, Bernardini+19, Chardin+19, Mustafa+19, Ramanah+20, Tamasiunas+20, Feder+20, Moster+20, Thiele+20, Wadekar+20, Dai+20, Li+20, Lucie-Smith+20, Kasmanoff+20, Ni+21, Rouhaiainen+21, Harrington+21, Horowitz+21, Horowitz+21, Bernardini+21, Schaurecker+21, Etezad-Razavi+21, Curtis+21


SEQUENCE MODELLING

*based on MIT Lecture by Ava Soleimany

GIVEN AN IMAGE OF A BALL, CAN YOU PREDICT
WHERE IT WILL GO NEXT?



GIVEN AN IMAGE OF A BALL, CAN YOU PREDICT
WHERE IT WILL GO NEXT?



Previous positions help guessing the future

LET'S TAKE A SIMPLE EXAMPLE OF LANGUAGE MODELLING

“This morning I took my cat for a walk.”

given these words

predict the
next word

LET'S TAKE A SIMPLE EXAMPLE OF LANGUAGE MODELLING


“This morning I took my cat for a walk.”

given these words

predict the
next word

Normal ANNs and CNNs cannot handle variable length inputs ...

RNNs for Sequence Modeling



WE COULD SIMPLY USE A FIXED WINDOW...


“This morning I took my cat for a walk.”

given these words

predict the
next word

[1 0 0 0 0 0 1 0 0 0]

for a




prediction

HOWEVER THIS CANNOT HANDLE LONG TERM MEMORY


“France is where I grew up, but I now live in Boston. I speak fluent ____.”

IN SOME CASES, INFORMATION FROM THE DISTANT PAST
IS NEEDED FOR A CORRECT PREDICTION



Moller+19

ANOTHER ALTERNATIVE COULD BE TO FEED A REALLY LARGE WINDOW




BUT WE STILL HAVE THE PROBLEM THAT THERE IS NO PARAMETER SHARING (SAME AS PIXELS)

RECURRENT NEURAL NETWORKS (RNNs)


OUR WISH LIST:

1. Handle variable-length sequences
2. Track long term dependencies
3. Maintain information about order
4. Share parameters across the sequence


Recurrent Neural Networks
offer a first solution to this problem




Handling Individual Time Steps




Handling Individual Time Steps




Neurons with Recurrence



Neurons with Recurrence




The RNN Block



$$h_t = f_W(h_{t-1}, x_t)$$

cell state function parameterized by W old state input vector at time step t

The RNN Block



Output Vector


$$\hat{y}_t = \mathbf{W}_{hy}^T h_t$$

Update Hidden State

$$h_t = \tanh(\mathbf{W}_{hh}^T h_{t-1} + \mathbf{W}_{xh}^T x_t)$$


Input Vector


$$x_t$$




RNNs: Computational Graph Across Time


Re-use the **same weight matrices** at every time step






BACKPROPAGATION THROUGH TIME (BPTT)





BPTT implies a large amount of weight multiplications: if we want to take into account long term memory, networks become quickly very deep and so are subjected to vanishing and exploding gradients problems (see previous lecture)




BPTT implies a large amount of weight multiplications: if we want to take into account long term memory, networks become quickly very deep and so are subjected to vanishing and exploding gradients problems (see previous lecture)

ONE CAN USE THE SAME TRICKS THAT WE DISCUSSED
FOR DEEP CNNs (WEIGHT INITIALISATION, RELU
ACTIVATION FUNCTION ...)


ANOTHER SPECIFIC SOLUTION: GATED CELLS

Use a more complex recurrent unit with gates to control what information is passed through...



ANOTHER SPECIFIC SOLUTION: GATED CELLS

Use a more complex recurrent unit with gates to control what information is passed through...



Long Short Term Memory (LSTMs) networks are an example of
this